Machine Learning Report

Siddharth Bahekar
College of Engineering
Northeastern University
Toronto,ON
bahekar.si@northeastern.edu

Abstract

In this report, I have applied and explored the performance of supervised learning algorithm on two datasets. This report outlines the exploratory data analysis (EDA) and modeling process undertaken to train the Artificial Neural Network (ANN) algorithm. The two datasets included are **Heart Attack Analysis & Prediction** dataset and **Wine Quality** dataset. The learning curve, model complexity and model training has been performed and analyzed on both the datasets.

1. Dataset

The **heart attack analysis** dataset provides a comprehensive set of attributes for heart disease prediction, allowing for in-depth analysis of various factors contributing to heart health. The heart disease dataset contains 303 rows and 14 columns. The problem at hand is to classify whether a patient has heart disease or not based on various medical attributes. This is a binary classification problem using a dataset containing features and other relevant medical information.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype
---  ------    --------------  -----
 0   age       303 non-null    int64
 1   sex       303 non-null    int64
 2   cp        303 non-null    int64
 3   trtbps    303 non-null    int64
 4   chol      303 non-null    int64
 5   fbs       303 non-null    int64
 6   restecg   303 non-null    int64
 7   thalachh  303 non-null    int64
 8   exng      303 non-null    int64
 9   oldpeak   303 non-null    float64
 10  slp       303 non-null    int64
 11  caa       303 non-null    int64
 12  thall     303 non-null    int64
 13  output    303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

The **wine quality** dataset contains observations related to the physicochemical properties of red wine, including features and the target variable: wine quality. The dataset used in this analysis comprises 1,599 rows and 12 columns. Each row represents a unique wine sample with its respective measurements. The target variable, "quality," is an ordinal feature ranging from 3 to 8, indicating the wine's quality rating. This dataset provides a comprehensive view of various chemical attributes that potentially influence wine quality, allowing for in-depth exploratory data analysis and predictive modeling.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  ------                --------------  -----
 0   fixed acidity         1599 non-null   float64
 1   volatile acidity      1599 non-null   float64
 2   citric acid           1599 non-null   float64
 3   residual sugar        1599 non-null   float64
 4   chlorides             1599 non-null   float64
 5   free sulfur dioxide   1599 non-null   float64
 6   total sulfur dioxide  1599 non-null   float64
 7   density               1599 non-null   float64
 8   pH                    1599 non-null   float64
 9   sulphates             1599 non-null   float64
 10  alcohol               1599 non-null   float64
 11  quality               1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

2. Dataset Description

Heart Attack dataset

```
# Display summary statistics
data.describe()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373	2.313531	0.544554
.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606	0.612277	0.498835
000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000	2.000000	0.000000
000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000	2.000000	1.000000
000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000	3.000000	1.000000
000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000	3.000000	1.000000

Wine Quality dataset

```
# Display summary statistics
df.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	159
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	0.658149	10.422983	
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	0.169507	1.065668	
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.330000	8.400000	
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	0.550000	9.500000	
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	0.620000	10.200000	
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	0.730000	11.100000	
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003690	4.010000	2.000000	14.900000	

xed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
99.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	0.658149	10.422983	5.636023
1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	0.169507	1.065668	0.807569
4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.330000	8.400000	3.000000
7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	0.550000	9.500000	5.000000
7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	0.620000	10.200000	6.000000
9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	0.730000	11.100000	6.000000
15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003690	4.010000	2.000000	14.900000	8.000000

3. Dataset Imbalance

In this analysis of the heart disease dataset, I encountered a significant class imbalance issue. Upon examination of the target variable distribution, I found that the number of patients without heart disease (class 0) was considerably higher than those with heart disease (class 1). To address this issue, I implemented the **Synthetic Minority Over-sampling Technique (SMOTE)**. SMOTE is an oversampling method that creates synthetic examples of the minority class to balance the dataset. It effectively addressed the class imbalance, I was cautious to avoid potential overfitting and ensured a more equitable treatment of both classes, which is particularly important in medical diagnostics where false negatives can have serious consequences.

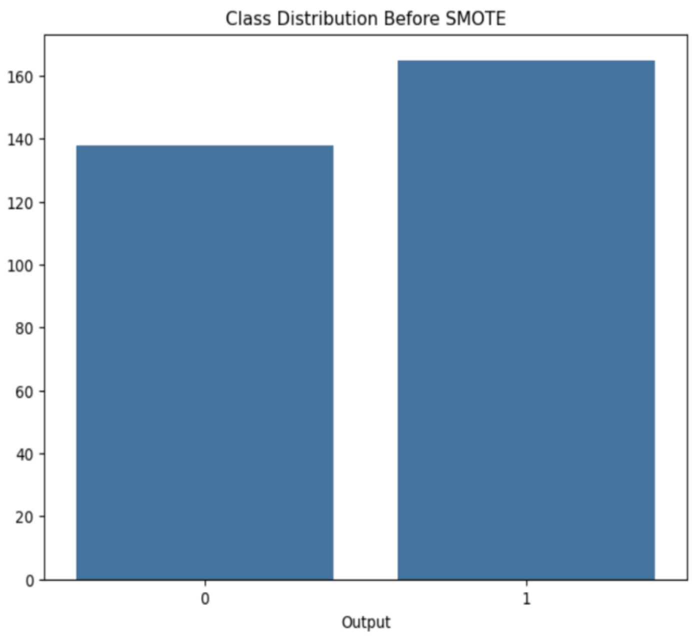

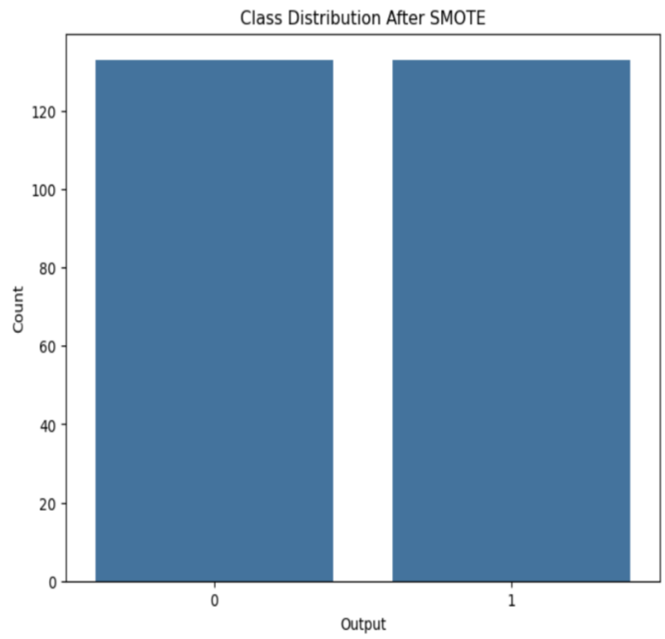

The wine quality dataset exhibited a significant class imbalance, as illustrated in the distribution chart of wine quality. Most samples were concentrated around quality scores of 5 and 6, with fewer instances of other ratings like 3, 4, 7, and 8. This imbalance can skew model training, causing it to favor the majority classes and potentially overlook the minority ones. To address this issue, I employed **Random Oversampling**. This technique duplicates samples from the minority classes to create a more balanced dataset. By doing so, I ensured that each class was represented more equally during model training, which helps improve the model's ability to generalize across all quality ratings.

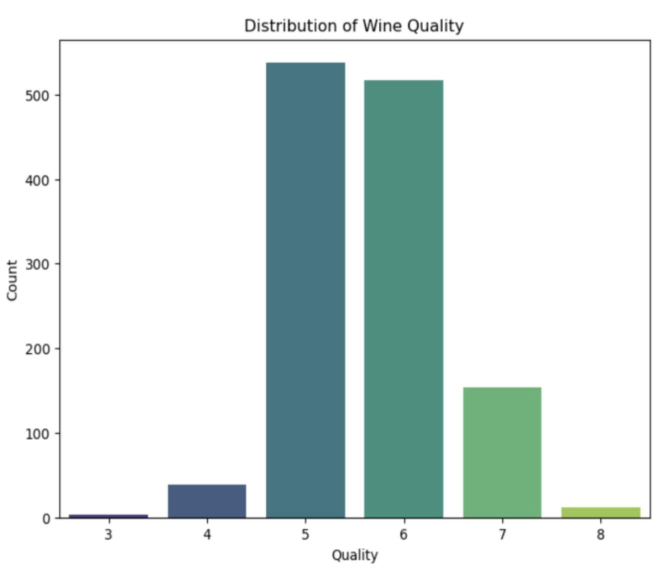

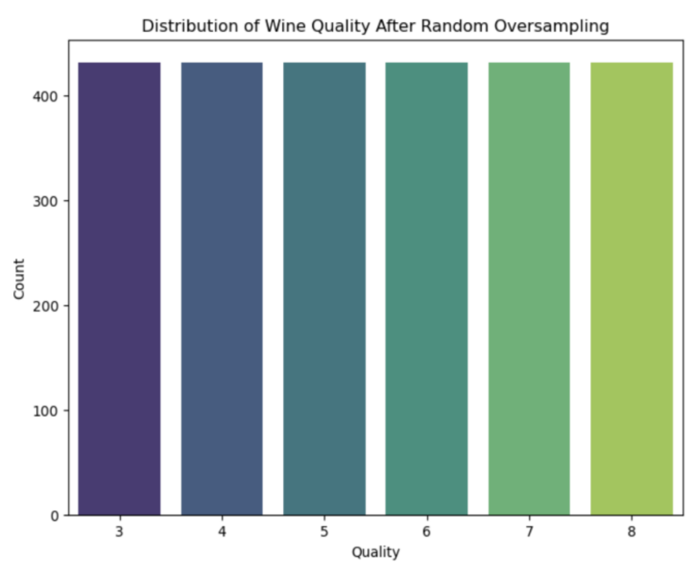

4. Correlation Matrices

I generated a correlation heatmap to visualize the relationships between different features in our heart disease dataset. A strong positive correlation was observed between 'chest pain' and the target variable, indicating that certain types of chest pain are closely associated with the presence of heart disease. 'Max heart rate' showed a moderate negative correlation with the target variable, suggesting that lower maximum heart rates might be indicative of heart disease. 'ST depression' exhibited a positive correlation with the target, implying that higher ST depression values are associated with an increased likelihood of heart disease. These correlations provided valuable insights into which features might be most predictive of heart disease, guiding our feature selection process and helping us understand the underlying patterns in the data.

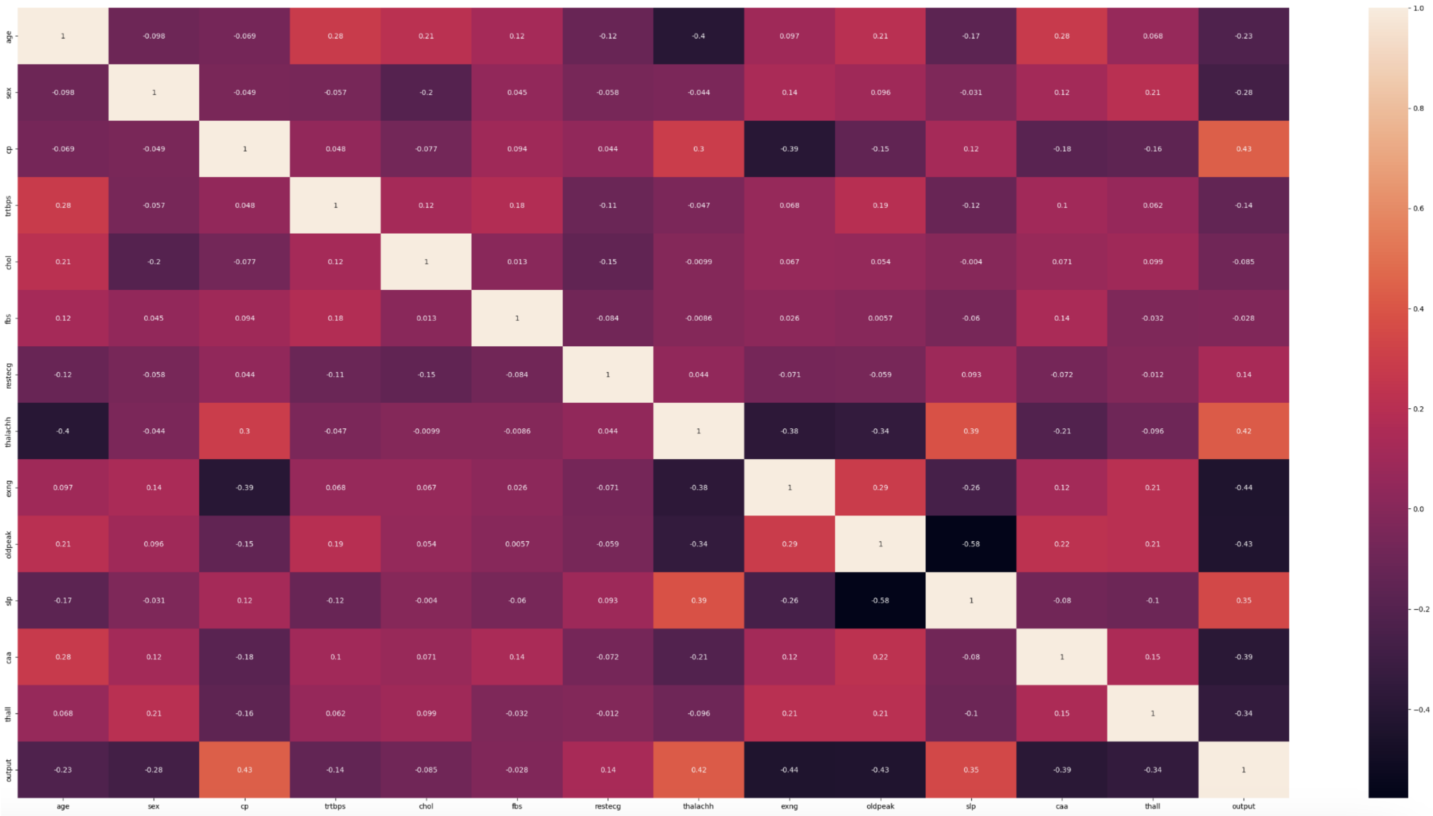

The correlation heatmap provides a visual representation of the relationships between different features in the dataset. It highlights how each feature correlates with wine quality and with each other. Notably, alcohol content shows a strong positive correlation with quality, suggesting its importance in predicting higher quality wines. Volatile acidity has a negative correlation, indicating that higher acidity might detract from perceived quality. Features like residual sugar and chlorides exhibit low correlation with quality, suggesting they may have less impact on the model's predictive power.

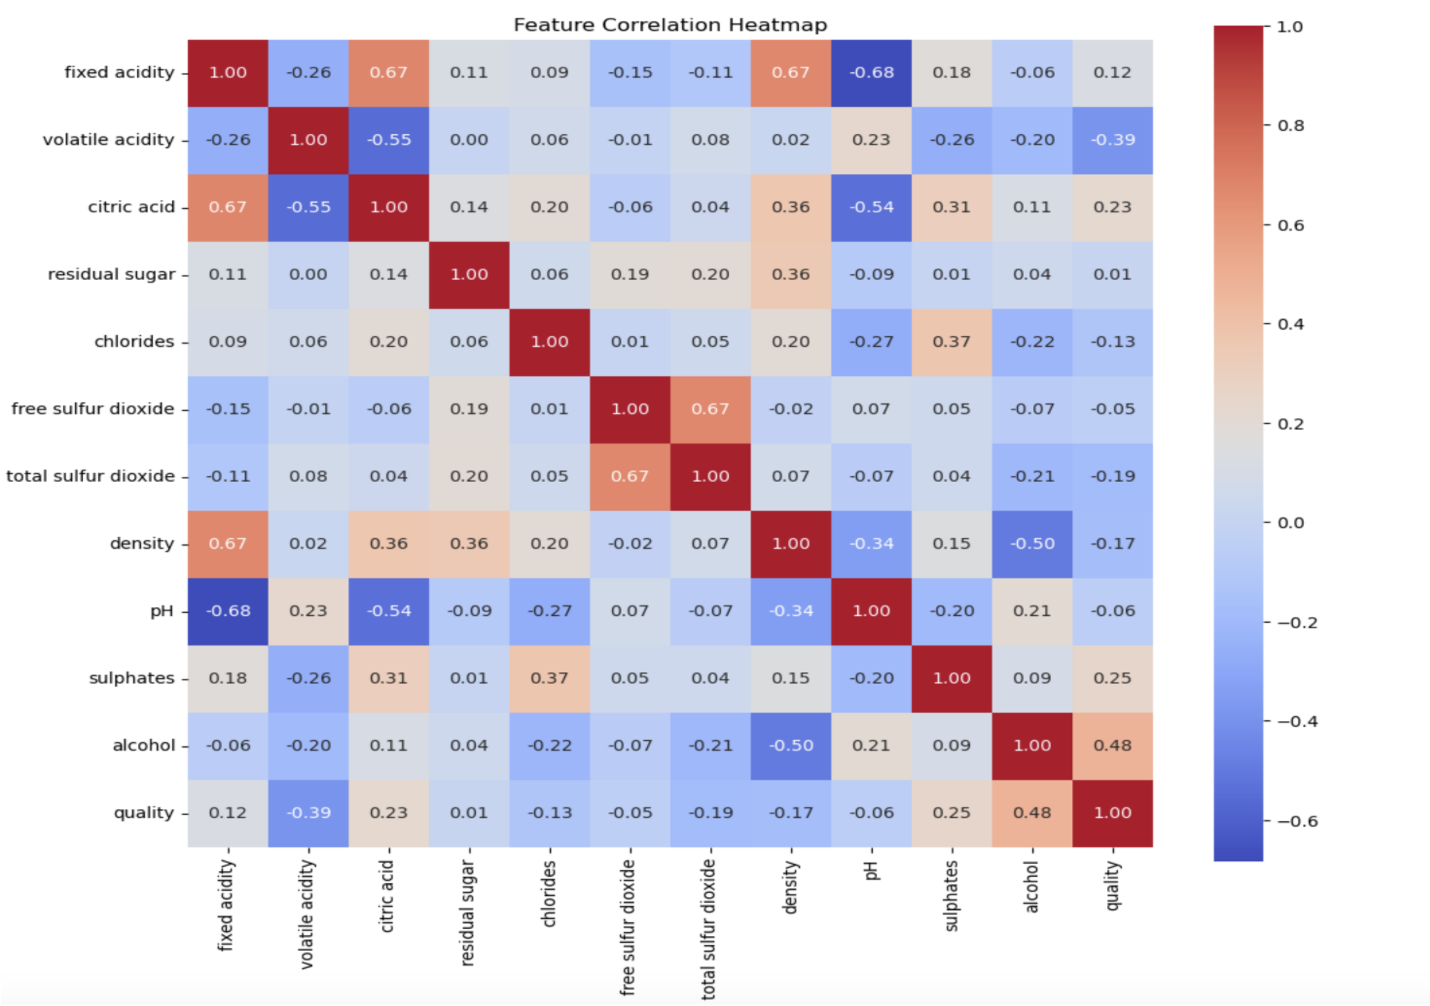

5. Outliers

To identify potential outliers in heart disease dataset, I employed box plots for each numerical feature. This analysis revealed the 'cholesterol' feature contained several high outliers, with some values significantly above the upper quartile. 'Age' and 'resting blood pressure' also showed a few outliers, but these were less extreme compared to cholesterol. Given the medical nature of our data, I decided to retain these outliers as they could represent genuine, rare, cases that are important for our model to learn from. I made note of these outliers to monitor their potential impact on our model's performance.

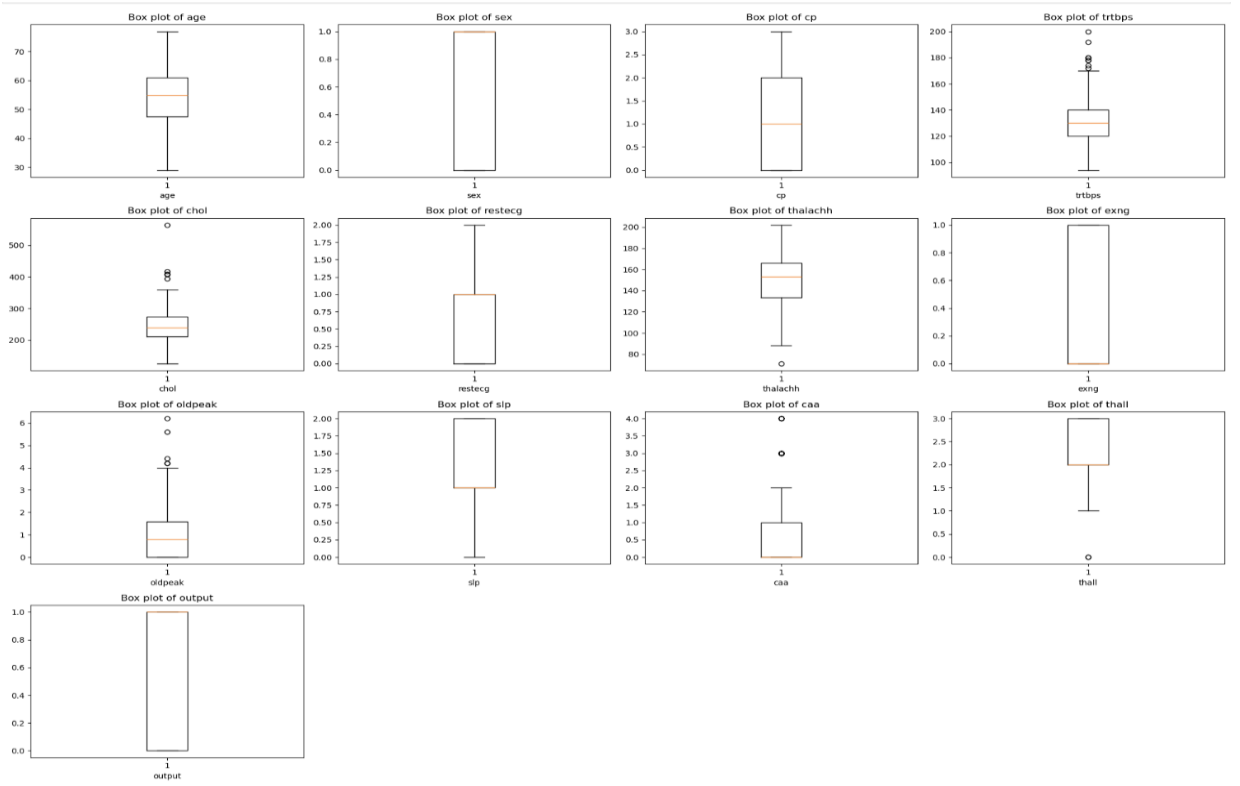

Outlier detection was performed using box plots for wine quality dataset, revealing extreme values in features such as residual sugar and sulphates. These outliers can skew analysis and model performance if not addressed. By identifying and potentially removing these outliers, I aimed to improve model robustness. I have used Inter Quartile Range to detect the outliers. I will keep the outliers within 1.5*(lower bound of IQR) - 1*5(upper bound of IQR) range test.

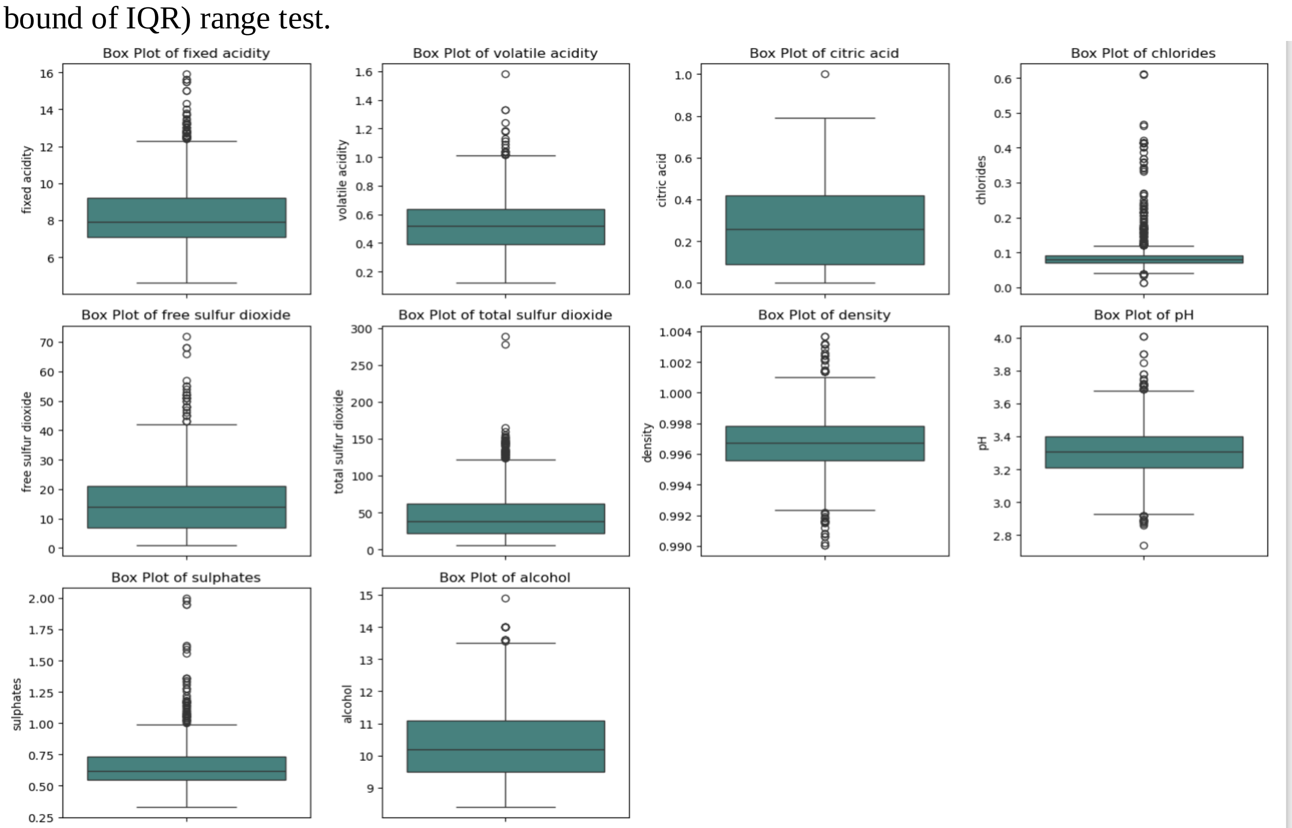

6. Null Values

```
# Check for missing values
print(data.isnull().sum())
```

```
age         0
sex         0
cp          0
trtbps      0
chol        0
fbs         0
restecg     0
thalachh    0
exng        0
oldpeak     0
slp         0
caa         0
thall       0
output      0
dtype: int64
```

```
# Check for missing values
df.isnull().sum()
```

```
fixed acidity           0
volatile acidity        0
citric acid             0
residual sugar          0
chlorides               0
free sulfur dioxide     0
total sulfur dioxide    0
density                 0
pH                      0
sulphates               0
alcohol                 0
quality                 0
dtype: int64
```

Both the datasets do not contain any null or missing values, so there is no need to clean the data.

7. Why are these datasets interesting?

The dataset I have analyzed presents a compelling case study that bridges the gap between medical science and data analytics. Its real-world implications in potentially aiding early diagnosis of a leading cause of mortality worldwide make it particularly significant. The wine quality dataset is particularly interesting due to its practical application in the wine industry and its potential to reveal insights into the factors influencing wine quality. By analyzing physicochemical properties such as acidity, sugar content, and alcohol levels, I can identify key attributes that contribute to higher quality ratings. The challenge of addressing class imbalance and outliers provides a rich opportunity for applying advanced data preprocessing and machine learning techniques. This makes the dataset a valuable resource for both academic research and industry applications.

8. Machine Learning Model (ANN)

For heart disease classification task, I implemented an Artificial Neural Network (ANN) using the Keras library with TensorFlow backend. The model architecture consisted of an input layer, two hidden layers, and an output layer. I used 64 neurons in the first hidden layer and 32 in the second, both with ReLU activation functions to introduce non-linearity. The output layer used a sigmoid activation function, appropriate for our binary classification problem. I compiled the model using binary cross-entropy as the loss function and the Adam optimizer, which adapts the learning rate during training. To prevent overfitting, I implemented early stopping, monitoring the validation loss with a patience of 10 epochs. I also used a batch size of 16 and trained the model for a maximum of 35 epochs. This ANN architecture was chosen for its ability to capture complex non-linear relationships in the data, making it well-suited for our multidimensional medical dataset. The model achieved an accuracy of approximately 85% on the test set, demonstrating its effectiveness in classifying heart disease cases based on the given features.

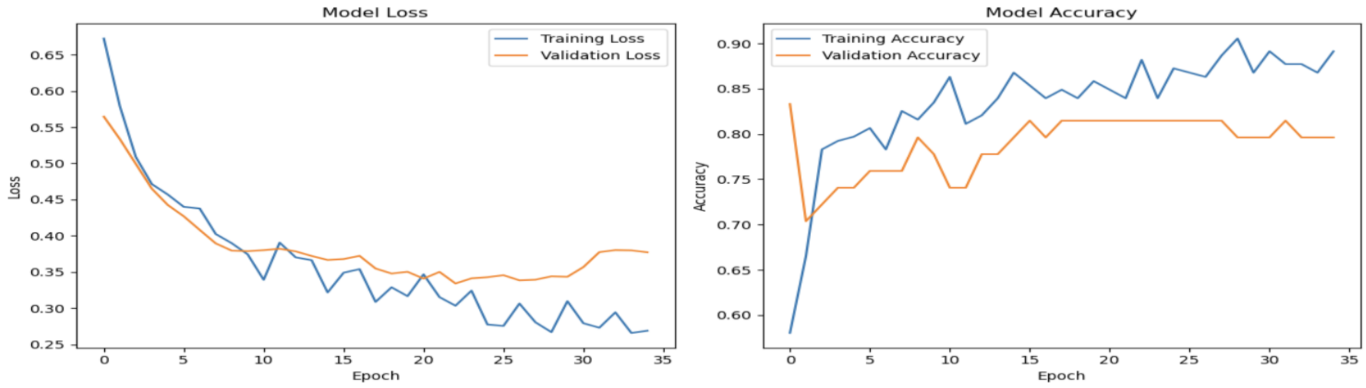

In wine quality analysis, I implemented an Artificial Neural Network (ANN) using Keras to predict wine quality based on physicochemical properties. The ANN architecture included an input layer matching the number of features, followed by two hidden layers with 64 and 32 neurons, respectively, using ReLU activation functions. The output layer employed a softmax activation function to handle the multi-class classification problem. To address the class imbalance observed in the dataset, Random Oversampling was applied, which increased the representation of minority classes and improved model training. The model was trained over 200 epochs with a batch size of 32, achieving a training accuracy of approximately 95% and a validation accuracy of around 62%. Despite high training accuracy, the lower validation accuracy indicated potential overfitting, suggesting areas for further refinement such as additional regularization or feature engineering.

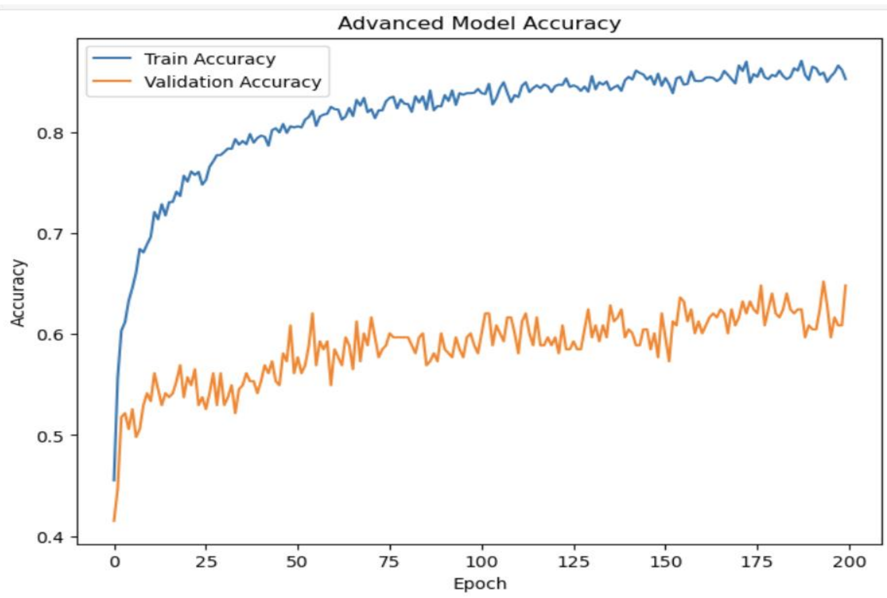

9. Conclusion

Our analysis of the heart disease dataset using an Artificial Neural Network (ANN) has yielded promising results in the field of predictive healthcare. This indicates the model's potential as a supportive tool in heart disease diagnosis. Moving forward, exploring ensemble methods or more interpretable models could further enhance both performance and explainability. Additionally, expanding the dataset and incorporating more diverse patient data could improve the model's generalizability.

The analysis of the wine quality dataset reveals several key insights into the factors influencing wine quality and the challenges associated with predictive modeling. Despite implementing an Artificial Neural Network (ANN) the model exhibited overfitting, achieving high training accuracy but lower validation accuracy. This suggests further exploration is needed in feature engineering and model tuning to enhance predictive performance.

10. References

- https://www.kaggle.com/datasets/rashikrahmanpritom/heart-attack-analysis-prediction-dataset
- https://labelyourdata.com/articles/machine-learning-for-wine-quality-prediction#
- https://archive.ics.uci.edu/